

# ADVANCED BAYESIAN MODELING

# Shakespeare Plays: JAGS Analysis

We use JAGS to fit the Shakespeare plays model – for illustration only.

The classical (non-Bayes) linear regression analysis would be adequate in this case – essentially equivalent to the noninformative prior model.

Start with classical analysis, since its estimates can be used to help initialize and make sure priors are not over-constrained ...

# Initial Classical Analysis

```
> mod <- lm(Speeches ~ 0 + Genre + scale(Words), data=ss)
```

Note:

- ▶ “0 +” prevents an explicit intercept from being added
- ▶ Genre is automatically coded using indicators, since it is a factor
- ▶ “scale” standardizes its argument

```
> summary(mod)
```

```
...
```

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )	
GenreComedy	1004.75	41.18	24.397	< 2e-16	***
GenreHistory	799.31	43.42	18.409	< 2e-16	***
GenreTragedy	1003.01	42.00	23.880	< 2e-16	***
scale(Words)	156.89	26.98	5.815	1.66e-06	***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 140.4 on 33 degrees of freedom
```

```
Multiple R-squared:  0.9811, Adjusted R-squared:  0.9789
```

```
F-statistic: 429.2 on 4 and 33 DF,  p-value: < 2.2e-16
```

All  $\beta$  estimates are about 100 to 1000, with much smaller standard errors.

We choose to set

- ▶ Initial  $\beta$  values at  $\pm 10000$

- ▶ Prior  $\beta$  variance at

$$\sigma_{\beta}^2 = (10000)^2$$

Regression error variance ( $\sigma^2$ ) estimate is about  $(140)^2 \approx 20000$ .

We'll choose initial  $\sigma^2$  values of 100 and 1000000.

In JAGS we use the equivalent prior (BDA3, Table A.1)

$$(\sigma^2)^{-1} \mid X \sim \text{Gamma}(\nu_0/2, \nu_0\sigma_0^2/2)$$

which is less informative as  $\nu_0 \rightarrow 0$ .

Arbitrarily, we set

$$\nu_0 = 0.0002 \qquad \sigma_0^2 = 1$$

# JAGS Analysis 1

```
model {  
  
  for (i in 1:length(y)) {  
    y[i] ~ dnorm(beta1*x1[i] + beta2*x2[i] + beta3*x3[i] + beta4*x4[i],  
                  sigmasqinv)  
  }  
  
  beta1 ~ dnorm(0, 0.00000001)  
  beta2 ~ dnorm(0, 0.00000001)  
  beta3 ~ dnorm(0, 0.00000001)  
  beta4 ~ dnorm(0, 0.00000001)  
  sigmasqinv ~ dgamma(0.0001, 0.0001)  
  
  sigmasq <- 1/sigmasqinv  
  
}
```



The data (in data frame `ss`) must be converted to appropriate numerical variables:

```
> d1 <- list(y = ss$Speeches,  
+           x1 = as.numeric(ss$Genre == "Comedy"),  
+           x2 = as.numeric(ss$Genre == "History"),  
+           x3 = as.numeric(ss$Genre == "Tragedy"),  
+           x4 = as.vector(scale(ss$Words)))
```

Set up initial values for 4 chains:

```
> inits1 <- list(list(beta1=10000, beta2=10000, beta3=10000, beta4=10000,  
+                    sigmasqinv=0.01),  
+               list(beta1=10000, beta2=10000, beta3=-10000, beta4=-10000,  
+                    sigmasqinv=0.01),  
+               list(beta1=10000, beta2=-10000, beta3=10000, beta4=-10000,  
+                    sigmasqinv=0.000001),  
+               list(beta1=10000, beta2=-10000, beta3=-10000, beta4=10000,  
+                    sigmasqinv=0.000001))
```

Set up and run chains as usual:

```
> library(rjags)
...

> m1 <- jags.model("shakespeare1.bug", d1, inits1, n.chains=4, n.adapt=1000)
...

> update(m1, 1000) # burn-in
|*****| 100%

> x1 <- coda.samples(m1, c("beta1","beta2","beta3","beta4","sigmasq"),
+                    n.iter=2000)
|*****| 100%
```

```
> gelman.diag(x1, autoburnin=FALSE)
```

Potential scale reduction factors:

	Point est.	Upper C.I.
beta1	1	1
beta2	1	1
beta3	1	1
beta4	1	1
sigmasq	1	1

Multivariate psrf

1

Graphical diagnostics (not shown) also show no problems with chains.

```
> effectiveSize(x1)
      beta1      beta2      beta3      beta4  sigmasq
5263.724 7018.103 6950.444 4451.034 6547.985
```

Sample sizes more than adequate.

```
> summary(x1)
```

```
Iterations = 1001:3000
```

```
Thinning interval = 1
```

```
Number of chains = 4
```

```
Sample size per chain = 2000
```

1. Empirical mean and standard deviation for each variable,  
plus standard error of the mean:

	Mean	SD	Naive SE	Time-series SE
beta1	1005.5	42.17	0.4715	0.5822
beta2	800.2	45.00	0.5031	0.5389
beta3	1002.2	43.55	0.4869	0.5260
beta4	157.6	27.91	0.3120	0.4265
sigmasq	21056.6	5520.99	61.7265	68.2147

(Compare with classical results – very similar.)

## 2. Quantiles for each variable:

	2.5%	25%	50%	75%	97.5%
beta1	921.6	977.2	1005.6	1033.4	1087.6
beta2	711.2	770.4	799.9	829.8	889.3
beta3	916.6	973.3	1002.5	1030.6	1088.2
beta4	102.8	139.3	157.5	176.1	212.7
sigmasq	12877.2	17150.2	20174.2	24067.9	34366.2

How probable is  $\beta_1 < \beta_3$ ? How about  $\beta_2 < \beta_1$ ?

```
> post.samp <- as.matrix(x1)

> mean(post.samp[, "beta1"] < post.samp[, "beta3"])
[1] 0.486
> # do comedies have fewer speeches than tragedies, adjusting for words?

> mean(post.samp[, "beta2"] < post.samp[, "beta1"])
[1] 0.999125
> # do histories have fewer speeches than comedies, adjusting for words?
```

So, after adjusting for words,

- ▶ No apparent evidence that comedies have fewer speeches than tragedies
- ▶ Strong evidence that histories have fewer speeches than comedies



The once-disputed comedy *The Two Noble Kinsmen* (not in the data set) has 25424 words.

A 95% central posterior interval for the *mean* number of speeches we expect it to have:

```
> quantile(post.samp[, "beta1"] +  
+          post.samp[, "beta4"] * (25424 - mean(ss$Words))/sd(ss$Words),  
+          c(0.025, 0.975))  
      2.5%      97.5%  
1017.114 1230.824
```

Note: For the *actual* number of speeches, need a posterior *predictive* interval.

## JAGS Analysis 2

We can use an alternative version of the JAGS model that generalizes more easily ...

```
model {  
  
  for (i in 1:length(y)) {  
    y[i] ~ dnorm(inprod(X[i,], beta), sigmasqinv)  
  }  
  
  beta ~ dmnorm(beta0, Sigmabetainv)  
  sigmasqinv ~ dgamma(0.0001, 0.0001)  
  
  sigmasq <- 1/sigmasqinv  
  
}
```

- ▶ `inprod` is the (vector) inner product
- ▶ `dmnorm` is the multivariate normal distribution, in a parameterization using the **precision matrix** (inverse covariance matrix)
- ▶ `beta0` and `Sigmabetainv` are specified with the data

Now the data (and constants in the priors) can be set up as

```
> d2 <- list(y = ss$Speeches,  
+           X = model.matrix(~ 0 + Genre + scale(Words), data=ss),  
+           beta0 = c(0,0,0,0),  
+           Sigmabetainv = rbind(c(0.00000001, 0, 0, 0),  
+                                c(0, 0.00000001, 0, 0),  
+                                c(0, 0, 0.00000001, 0),  
+                                c(0, 0, 0, 0.00000001)))
```

Initialization of beta is also modified:

```
> inits2 <- list(list(beta=c(10000, 10000, 10000, 10000),  
+                      sigmasqinv=0.01),  
+               list(beta=c(10000, 10000, -10000, -10000),  
+                      sigmasqinv=0.01),  
+               list(beta=c(10000, -10000, 10000, -10000),  
+                      sigmasqinv=0.000001),  
+               list(beta=c(10000, -10000, -10000, 10000),  
+                      sigmasqinv=0.000001))
```

Results are almost exactly the same as before, except the Gibbs sampler runs a bit more efficiently (because of block updating).